

# Ftrace Kernel Hooks: More than just tracing

Presenter:

Steven Rostedt

[rostedt@goodmis.org](mailto:rostedt@goodmis.org)

Red Hat



[https://download.csdn.net/download/Rong\\_Toa/18654864](https://download.csdn.net/download/Rong_Toa/18654864)

# Ftrace Function Hooks

- Function Tracer
- Function Graph Tracer
- Function Profiler
- Stack Tracer
- Kprobes
- Uprobes
- Perf
- Pstore
- SystemTap

# Function Tracing

```
# cd /sys/kernel/debug/tracing
# echo function > current_tracer
# cat trace
# tracer: function
#
# entries-in-buffer/entries-written: 205022/119956607   #P:4
#
#          _-----=> irqsoff
#          /_-----=> need-resched
#          | /_-----=> hardirq/softirq
#          || /_-----=> preempt-depth
#          ||| /_-----=> delay
#
# TASK-PID   CPU#  | TIMESTAMP | FUNCTION
# | |       |   |   |   |   |
<idle>-0    [002] dN.1  1781.978299: rcu_eqs_exit <-rcu_idle_exit
<idle>-0    [002] dN.1  1781.978300: rcu_eqs_exit_common <-rcu_eqs_exit
<idle>-0    [002] .N.1  1781.978301: arch_cpu_idle_exit <-cpu_startup_entry
<idle>-0    [002] .N.1  1781.978301: tick_nohz_idle_exit <-cpu_startup_entry
<idle>-0    [002] dN.1  1781.978301: ktime_get <-tick_nohz_idle_exit
<idle>-0    [002] dN.1  1781.978302: update_ts_time_stats <-tick_nohz_idle_exit
<idle>-0    [002] dN.1  1781.978302: nr_iowait_cpu <-update_ts_time_stats
<idle>-0    [002] dN.1  1781.978303: tick_do_update_jiffies64 <-tick_nohz_idle_exit
<idle>-0    [002] dN.1  1781.978303: update_cpu_load_nohz <-tick_nohz_idle_exit
<idle>-0    [002] dN.1  1781.978303: calc_load_exit_idle <-tick_nohz_idle_exit
```

# Function Graph Tracer

```
# echo function_graph > current_tracer
# cat trace
# tracer: function_graph
#
# CPU    DURATION    FUNCTION CALLS
# |      |      |      |      |      |
2)      7.879 us    | } /* context_tracking_user_exit */
2)      |          | __do_page_fault() {
2)      0.070 us    | down_read_trylock();
2)      0.057 us    | __might_sleep();
2)      0.096 us    | find_vma();
2)      |          | handle_mm_fault() {
2)      |          |     __do_fault() {
2)      |          |         filemap_fault() {
2)      |          |             find_get_page() {
2)      |          |                 __rcu_read_lock();
2)      0.057 us    |                 __rcu_read_unlock();
2)      0.061 us    |             }
2)      1.241 us    |             __might_sleep();
2)      0.074 us    |         }
2)      2.201 us    |     }
2)      |          | _raw_spin_lock() {
2)      0.069 us    |     preempt_count_add();
2)      0.528 us    | }
2)      0.063 us    | add_mm_counter_fast();
2)      0.070 us    | page_add_file_rmap();
2)      |          | _raw_spin_unlock() {
2)      0.070 us    |     preempt_count_sub();
```

# Dynamic Function Tracing

```
# echo '*sched*' > set_ftrace_filter
# echo function > current_tracer
# cat trace
# tracer: function
#
# entries-in-buffer/entries-written: 193727/240417   #P:4
#
#          _-----=> irqs-off
#          /_-----=> need-resched
#          |/_-----=> hardirq/softirq
#          ||/_-----=> preempt-depth
#          |||/_-----=> delay
#
# TASK-PID   CPU#  | TIMESTAMP | FUNCTION
# | |         |   |         |   |
<idle>-0    [003] d.h3      6325.742705: resched_task <-check_preempt_curr
<idle>-0    [003] dNh3      6325.742712: native_smp_send_reschedule <-enqueue_task_fair
<idle>-0    [003] dNh3      6325.742714: resched_task <-check_preempt_curr
<idle>-0    [003] dN.1      6325.742719: smp_reschedule_interrupt <-reschedule_interrupt
<idle>-0    [003] dN.1      6325.742720: scheduler_ipi <-smp_reschedule_interrupt
<idle>-0    [003] dNh1      6325.742722: sched_ttwu_pending <-scheduler_ipi
<idle>-0    [003] .N.1      6325.742728: schedule_preempt_disabled <-cpu_startup_entry
<idle>-0    [003] .N..      6325.742729: schedule <-schedule_preempt_disabled
<idle>-0    [003] .N..      6325.742731: __schedule <-preempt_schedule
<idle>-0    [003] .N.1      6325.742732: rcu_sched_qs <-rcu_note_context_switch
<idle>-0    [003] dN.2      6325.742733: pre_schedule_idle <-__schedule
aprsd-3467  [003] ....      6325.742746: schedule <-do_nanosleep
aprsd-3467  [003] ....      6325.742747: __schedule <-schedule
aprsd-3467  [003] ...1      6325.742748: rcu_sched_qs <-rcu_note_context_switch
aprsd-3454  [003] ....      6325.742767: schedule <-do_nanosleep
aprsd-3454  [003] ....      6325.742767: __schedule <-schedule
aprsd-3454  [003] ...1      6325.742768: rcu_sched_qs <-rcu_note_context_switch
rcu_preempt-9 [003] d..2      6325.742788: smp_reschedule_interrupt <-reschedule_interrupt
rcu_preempt-9 [003] d..2      6325.742789: scheduler_ipi <-smp_reschedule_interrupt
```

# How it works?

- gcc's profiler option: **-pg**
- Adds special mcount function call
  - all functions call mcount
  - mcount is a trampoline 蹦床



# A function call

```
asmlinkage __visible void __sched schedule(void)
{
    struct task_struct *tsk = current;

    sched_submit_work(tsk);
    __schedule();
}
```

# A function call

- Disassembled

<schedule>:

55		push	%rbp
48	8b 04 25 80 c0 0e	mov	0xffffffff810ec080,%rax
81			
48	89 e5	mov	%rsp,%rbp
48	8b 00	mov	(%rax),%rax
5d		pop	%rbp
e9	db fa ff ff	jmpq	ffffffff810bb100 <__schedule>
66	66 2e 0f 1f 84 00	data16	nopw %cs:0x0(%rax,%rax,1)
00	00 00 00 00		



# A function call With -pg option


- Disassembled

```
<schedule>:
  55                                push    %rbp
  48 89 e5                          mov     %rsp,%rbp
  e8 37 2e 00 00                   callq   ffffffff810f7430 <mcount>
  5d                                pop     %rbp
  48 8b 04 25 80 d0 15             mov     0xffffffff8115d080,%rax
  81
  48 8b 00                          mov     (%rax),%rax
  e9 96 fa ff ff                  jmpq    ffffffff810f40a0 <__schedule>
  66 0f 1f 44 00 00               nopw    0x0(%rax,%rax,1)
```

# The kernel at boot up

```
<schedule>:  
push    %rbp  
mov     %rsp,%rbp  
callq   <mcount>  
pop     %rbp
```

```
<mcount>:  
retq
```



# Where's the mcounts?

- Can't just call the mcounts
  - too much overhead
  - retq only is 13% added overhead!
- Need to convert them to nops at boot up  
启动过程均为nop指令
- Need to know where they are  
他们在哪
- Best to find them at compile time  
最好是在编译过程就发现他们

# recordmcount

- scripts/recordmcount.c (and a perl version)
- reads the object files one at a time
- reads the relocation tables
  - finds all the calls to mcount
  - creates a table
  - links the table back into the object file
  - New section called **\_\_mcount\_loc**

参见内核中的section

```
__start_mcount_loc = .;          /* ffffffff83dfa990 */ \
KEEP(*(__mcount_loc))           \
KEEP(*(__patchable_function_entries)) \
__stop_mcount_loc = .;          /* ffffffff83e4a800 */ \
```

# recordmcount

kernel/sched/core.o:

```
<schedule>:  
push    %rbp  
mov     %rsp,%rbp  
callq  <mcount>  
pop     %rbp  
[...]  
<preempt_schedule_irq>:  
push    %rbp  
mov     %rsp,%rbp  
push    %rbx  
callq  <mcount>  
pop     %rbp  
[...]  
<_cond_resched>:  
push    %rbp  
mov     %rsp,%rbp  
push    %rbx  
callq  <mcount>  
pop     %rbp  
[...]  
<yield>:  
push    %rbp  
mov     %rsp,%rbp  
push    %rbx  
callq  <mcount>  
pop     %rbp  
[...]
```

这就是正经八百的牛逼

```
<__mcount_loc>:  
&schedule + 0x4  
&preempt_schedule_irq + 0x4  
&_cond_resched + 0x4  
&yield + 0x4
```

# recordmcount

kernel/sched/core.o:

```
<schedule>:
push    %rbp
mov     %rsp,%rbp
callq  <mcount>
pop     %rbp
[...]
<preempt_schedule_irq>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
callq  <mcount>
pop     %rbp
[...]
<_cond_resched>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
callq  <mcount>
pop     %rbp
[...]
<yield>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
callq  <mcount>
pop     %rbp
[...]
<__mcount_loc>:
&schedule + 0x4
&preempt_schedule_irq + 0x4
&_cond_resched + 0x4
&yield + 0x4
```

将其添加到.o文件中



# Linker Magic

- vmlinux.lds
  - include/linux/vmlinux.lds.h
  - arch/x86/kernel/vmlinux.lds.S
- Magic variables
  - \_\_start\_mcount\_loc
  - \_\_stop\_mcount\_loc

```
#ifdef CONFIG_FTRACE_MCOUNT_RECORD
#define MCOUNT_REC()      . = ALIGN(8);                \
                        VMLINUX_SYMBOL(__start_mcount_loc) = .; \
                        *(__mcount_loc)                    \
                        VMLINUX_SYMBOL(__stop_mcount_loc) = .;

#else
#define MCOUNT_REC()
#endif
```

# Linker Magic

vmlinux:



kernel/sched/core.o:

```
<__mcount_loc>:  
&schedule + 0x4  
&preempt_schedule_irq + 0x4  
&_cond_resched + 0x4  
&yield + 0x4
```

mm/swap.o:

```
<__mcount_loc>:  
&put_page + 0x4  
&__get_page_tail + 0x4  
&put_pages_list + 0x4  
&get_kernel_pages + 0x4
```

fs/read\_write.o:

```
<__mcount_loc>:  
&new_sync_read + 0x4  
&vfs_setpos + 0x4  
&fixed_size_llseek + 0x4  
&default_llseek + 0x4
```

...

# Linker Magic

vmlinux:

```
<__start_mcount_loc>:  
&schedule + 0x4  
&preempt_schedule_irq + 0x4  
&cond_resched + 0x4  
&yield + 0x4  
&put_page + 0x4  
&__get_page_tail + 0x4  
&put_pages_list + 0x4  
&get_kernel_pages + 0x4  
&new_sync_read + 0x4  
&vfs_setpos + 0x4  
&fixed_size_llseek + 0x4  
&default_llseek + 0x4  
[...]  
<__end_mcount_loc>:
```

kernel/sched/core.o:

```
<__mcount_loc>:  
&schedule + 0x4  
&preempt_schedule_irq + 0x4  
&cond_resched + 0x4  
&yield + 0x4
```

mm/swap.o:

```
<__mcount_loc>:  
&put_page + 0x4  
&__get_page_tail + 0x4  
&put_pages_list + 0x4  
&get_kernel_pages + 0x4
```

fs/read\_write.o:

```
<__mcount_loc>:  
&new_sync_read + 0x4  
&vfs_setpos + 0x4  
&fixed_size_llseek + 0x4  
&default_llseek + 0x4
```

...

# Linker Magic

vmlinux:

<\_\_start\_mcount\_loc>:

```
0xffffffff810f45f4
0xffffffff810f4635
0xffffffff810f4684
0xffffffff810f4734
0xffffffff81087ad4
0xffffffff81087b14
0xffffffff81087bd5
0xffffffff81087c41
0xffffffff810a7aa0
0xffffffff810a7bd4
0xffffffff810a7d34
0xffffffff810a7d7d
```

[...]

<\_\_end\_mcount\_loc>:

kernel/sched/core.o:

```
<__mcount_loc>:
&schedule + 0x4
&preempt_schedule_irq + 0x4
&_cond_resched + 0x4
&yield + 0x4
```

mm/swap.o:

```
<__mcount_loc>:
&put_page + 0x4
&__get_page_tail + 0x4
&put_pages_list + 0x4
&get_kernel_pages + 0x4
```

fs/read\_write.o:

```
<__mcount_loc>:
&new_sync_read + 0x4
&vfs_setpos + 0x4
&fixed_size_llseek + 0x4
&default_llseek + 0x4
```

...

# Finding Mcount

vmlinux:

```
<schedule>:
push    %rbp
mov     %rsp,%rbp
callq  <mcount>
pop     %rbp
[...]
<preempt_schedule_irq>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
callq  <mcount>
pop     %rbp
[...]
<_cond_resched>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
callq  <mcount>
pop     %rbp
[...]
<yield>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
callq  <mcount>
pop     %rbp
[...]

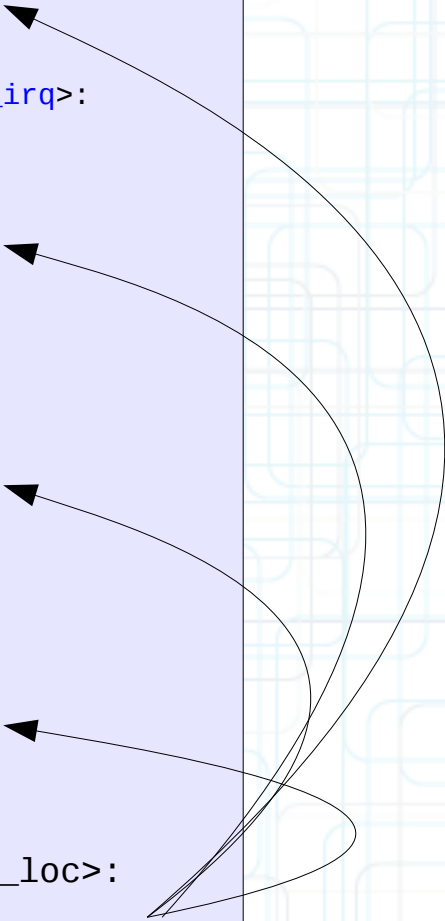
<__start_mcount_loc>:
[...]
<__end_mcount_loc>:
```

# Finding Mcount

vmlinux: 当开启mcount/ftrace功能

```
<schedule>:
push    %rbp
mov     %rsp,%rbp
callq  <mcount>
pop     %rbp
[...]
<preempt_schedule_irq>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
callq  <mcount>
pop     %rbp
[...]
<_cond_resched>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
callq  <mcount>
pop     %rbp
[...]
<yield>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
callq  <mcount>
pop     %rbp
[...]

<__start_mcount_loc>:
[...]
<__end_mcount_loc>:
```



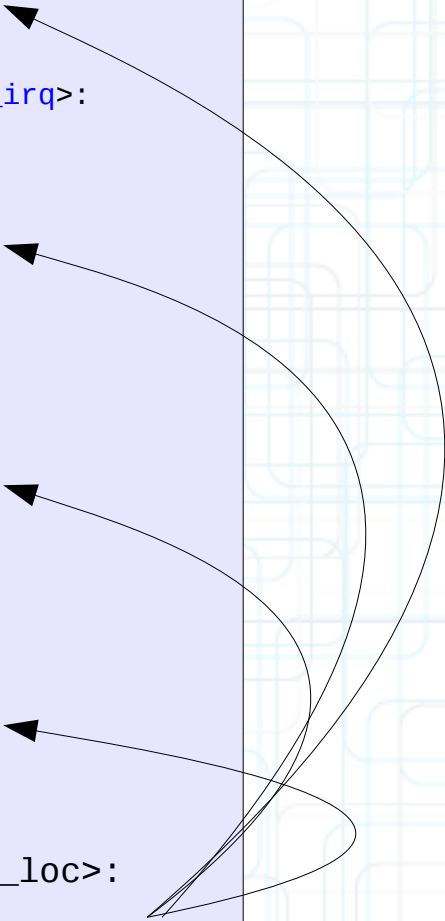


# Finding Mcount

vmlinux: 当关闭ftrace功能

```
<schedule>:
push    %rbp
mov     %rsp,%rbp
nop
pop     %rbp
[...]
<preempt_schedule_irq>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
nop
pop     %rbp
[...]
<_cond_resched>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
nop
pop     %rbp
[...]
<yield>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
nop
pop     %rbp
[...]

<__start_mcount_loc>:
[...]
<__end_mcount_loc>:
```



# Finding Mcount

vmlinux:

```
<schedule>:
push    %rbp
mov     %rsp,%rbp
nop
pop     %rbp
[...]
<preempt_schedule_irq>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
nop
pop     %rbp
[...]
<_cond_resched>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
nop
pop     %rbp
[...]
<yield>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
nop
pop     %rbp
[...]

<__start_mcount_loc>:
[...]
<__end_mcount_loc>:
```



# Finding Mcount

vmlinux:

```
<schedule>:  
push    %rbp  
mov     %rsp,%rbp  
nop  
pop     %rbp  
[...]  
<preempt_schedule_irq>:  
push    %rbp  
mov     %rsp,%rbp  
push    %rbx  
nop  
pop     %rbp  
[...]  
<_cond_resched>:  
push    %rbp  
mov     %rsp,%rbp  
push    %rbx  
nop  
pop     %rbp  
[...]  
<yield>:  
push    %rbp  
mov     %rsp,%rbp  
push    %rbx  
nop  
pop     %rbp  
[...]
```



# What about tracing?

使能方法

- Need a way to enable tracing
- We threw away the mcount section  
扔掉mcount section
- The mcount section wasn't enough for us
- Tracing also requires saving state

# struct dyn\_ftrace

```
struct dyn_ftrace {  
    unsigned long      ip; /* address of mcount call-site */  
    unsigned long      flags;  
    struct dyn_arch_ftrace arch;  
};
```

# struct dyn\_ftrace

```
struct dyn_ftrace {  
    unsigned long    ip; /* address of mcount call-site */  
    unsigned long    flags;  
    struct dyn_arch_ftrace arch;  
};
```

[arch/x86/include/asm/ftrace.h](#):

```
struct dyn_arch_ftrace {  
    /* No extra data needed for x86 */  
};
```



# struct dyn\_ftrace

```
struct dyn_ftrace {  
    unsigned long      ip; /* address of mcount call-site */  
    unsigned long      flags;  
    struct dyn_arch_ftrace arch;  
};
```

[arch/powerpc/include/asm/ftrace.h](#):

```
struct dyn_arch_ftrace {  
    struct module *mod;  
};
```

# Tracing data

- Copy from mcount\_loc before deleting
- Sorted for quick lookup 排序- 为了查找
- Allocated in groups of pages
  - details out of scope for this talk
- Data reported at boot up

```
# dmesg |grep ftrace
```

```
[ 0.139656] ftrace: allocating 24683 entries in 97 pages
```

- Allocated 24,683 dyn\_ftrace structures
- Used up 97 (4K) pages to do so
- Total of 397,312 bytes

# Tracing data

vmlinux:

```
<schedule>:
push    %rbp
mov     %rsp,%rbp
nop
pop     %rbp
[...]
<preempt_schedule_irq>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
nop
pop     %rbp
[...]
<_cond_resched>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
nop
pop     %rbp
[...]
<yield>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
nop
pop     %rbp
[...]

<__start_mcount_loc>:
[...]
<__end_mcount_loc>:
```

<ftrace\_pages>

```
ip      = 0xffffffff81087ad4
flags   = 0
ip      = 0xffffffff81087b14
flags   = 0
ip      = 0xffffffff81087bd5
flags   = 0
ip      = 0xffffffff81087c41
flags   = 0
ip      = 0xffffffff810a7aa0
flags   = 0
ip      = 0xffffffff810a7bd4
flags   = 0
ip      = 0xffffffff810a7d34
flags   = 0
ip      = 0xffffffff810a7d7d
flags   = 0
ip      = 0xffffffff810f45f4
flags   = 0
ip      = 0xffffffff810f4635
flags   = 0
ip      = 0xffffffff810f4684
flags   = 0
ip      = 0xffffffff810f4734
flags   = 0
[...]
```

# Tracing data

<trace\_pages> 可用的

```
ip    = 0xffffffff81087ad4
flags = 0
ip    = 0xffffffff81087b14
flags = 0
ip    = 0xffffffff81087bd5
flags = 0
ip    = 0xffffffff81087c41
flags = 0
ip    = 0xffffffff810a7aa0
flags = 0
ip    = 0xffffffff810a7bd4
flags = 0
ip    = 0xffffffff810a7d34
flags = 0
ip    = 0xffffffff810a7d7d
flags = 0
ip    = 0xffffffff810f45f4
flags = 0
ip    = 0xffffffff810f4635
flags = 0
ip    = 0xffffffff810f4684
flags = 0
ip    = 0xffffffff810f4734
flags = 0
[...]
```

# cat available\_filter\_functions

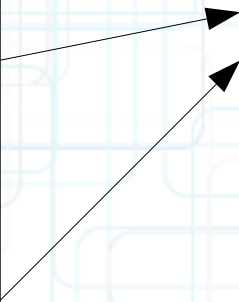
put\_page  
\_\_get\_page\_tail  
put\_pages\_list  
get\_kernel\_pages  
new\_sync\_read  
vfs\_setpos  
fixed\_size\_llseek  
default\_llseek  
schedule  
preempt\_schedule\_irq  
\_cond\_resched  
yield

# Tracing data

<ftrace\_pages>

```
ip      = 0xffffffff81087ad4
flags   = 0
ip      = 0xffffffff81087b14
flags   = 0
ip      = 0xffffffff81087bd5
flags   = 0
ip      = 0xffffffff81087c41
flags   = 0
ip      = 0xffffffff810a7aa0
flags   = 0
ip      = 0xffffffff810a7bd4
flags   = 0
ip      = 0xffffffff810a7d34
flags   = 0
ip      = 0xffffffff810a7d7d
flags   = 0
ip      = 0xffffffff810f45f4
flags   = 0
ip      = 0xffffffff810f4635
flags   = 0
ip      = 0xffffffff810f4684
flags   = 0
ip      = 0xffffffff810f4734
flags   = 0
[...]
```

```
# echo yield > set_ftrace_filter
# echo schedule >> set_ftrace_filter
# cat set_ftrace_filter
schedule
yield
```





# dyn\_ftrace.flags

- First 29 bits are for counter
  - Every registered callback increments +1
- bit 29 (starts from zero) – ENABLED
- bit 30 – REGS
- bit 31 – REGS\_EN



使能tracing

# Enabling tracing

vmlinux:

```
<schedule>:
push    %rbp
mov     %rsp,%rbp
nop
pop     %rbp
[...]
<preempt_schedule_irq>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
nop
pop     %rbp
[...]
<_cond_resched>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
nop
pop     %rbp
[...]
<yield>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
nop
pop     %rbp
[...]
```

<ftrace\_pages>

```
ip      = 0xffffffff81087ad4
flags   = 0
ip      = 0xffffffff81087b14
flags   = 0
ip      = 0xffffffff81087bd5
flags   = 0
ip      = 0xffffffff81087c41
flags   = 0
ip      = 0xffffffff810a7aa0
flags   = 0
ip      = 0xffffffff810a7bd4
flags   = 0
ip      = 0xffffffff810a7d34
flags   = 0
ip      = 0xffffffff810a7d7d
flags   = 0
ip      = 0xffffffff810f45f4
flags   = 0x20000001
ip      = 0xffffffff810f4635
flags   = 0
ip      = 0xffffffff810f4684
flags   = 0
ip      = 0xffffffff810f4734
flags   = 0xa0000001
[...]
```

# Enabling tracing

vmlinux:

```
<schedule>:
push    %rbp
mov     %rsp,%rbp
call ftrace_caller
pop     %rbp
[...]
<preempt_schedule_irq>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
nop
pop     %rbp
[...]
<_cond_resched>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
nop
pop     %rbp
[...]
<yield>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
call ftrace_regs_caller
pop     %rbp
[...]
```

<ftrace\_pages>

```
ip      = 0xffffffff81087ad4
flags   = 0
ip      = 0xffffffff81087b14
flags   = 0
ip      = 0xffffffff81087bd5
flags   = 0
ip      = 0xffffffff81087c41
flags   = 0
ip      = 0xffffffff810a7aa0
flags   = 0
ip      = 0xffffffff810a7bd4
flags   = 0
ip      = 0xffffffff810a7d34
flags   = 0
ip      = 0xffffffff810a7d7d
flags   = 0
ip      = 0xffffffff810f45f4
flags   = 0x20000001
ip      = 0xffffffff810f4635
flags   = 0
ip      = 0xffffffff810f4684
flags   = 0
ip      = 0xffffffff810f4734
flags   = 0xe0000001
[...]
```

# Modifying code at runtime

- Not the same as at boot up
- SMP boxes must be careful 多处理器需要注意  
其他的CPU可能执行这个代码
- Other CPUs may be executing that code  
x86具有非统一的机器指令
- x86 has non uniform machine instructions  
指令可能会跨越缓存边界
- Instructions may cross cache boundaries

# Modifying code at runtime

<schedule>:

55		push	%rbp
48 89 e5		mov	%rsp,%rbp
0f 1f 44 00 00		nop	
5d		pop	%rbp
48 8b 04 25 80 d0 15		mov	0xfffffffff8115d080,%rax
81			
48 8b 00		mov	(%rax),%rax
e9 96 fa ff ff		jmpq	fffffffff810f40a0 <__schedule>
66 0f 1f 44 00 00		nopw	0x0(%rax,%rax,1)

# Modifying code at runtime

<schedule>:

55		push	%rbp
48 89 e5		mov	%rsp,%rbp
e8 37 2e 00 00		callq	ffffffff810f7430 <ftrace_caller>
5d		pop	%rbp
48 8b 04 25 80 d0 15		mov	0xffffffff8115d080,%rax
81			
48 8b 00		mov	(%rax),%rax
e9 96 fa ff ff		jmpq	ffffffff810f40a0 <__schedule>
66 0f 1f 44 00 00		nopw	0x0(%rax,%rax,1)

# Modifying code at runtime

## CPU 0

<schedule>:

```
55
48 89 e5
0f 1f 44 00 00
5d
48 8b 04 25 80 d0 15
81
48 8b 00
e9 96 fa ff ff
66 0f 1f 44 00 00
```

## CPU 1

<schedule>:

```
→ 55
48 89 e5
0f 1f 44 00 00
5d
48 8b 04 25 80 d0 15
81
48 8b 00
e9 96 fa ff ff
66 0f 1f 44 00 00
```

0f 1f 44 00 00 nop

e8 37 2e 00 00 callq ffffffff810f7430 <ftrace\_caller>



# Modifying code at runtime

## CPU 0

```
<schedule>:  
55  
48 89 e5  
e8 37 2e 00 00  
5d  
48 8b 04 25 80 d0 15  
81  
48 8b 00  
e9 96 fa ff ff  
66 0f 1f 44 00 00
```

## CPU 1

```
<schedule>:  
55  
→ 48 89 e5  
0f 1f 44 00 00  
5d  
48 8b 04 25 80 d0 15  
81  
48 8b 00  
e9 96 fa ff ff  
66 0f 1f 44 00 00
```

```
0f 1f 44 00 00 nop  
e8 37 2e 00 00 callq ffffffff810f7430 <ftrace_caller>
```

# Modifying code at runtime

## CPU 0

```
<schedule>:
55
48 89 e5
e8 37 2e 00 00
5d
48 8b 04 25 80 d0 15
81
48 8b 00
e9 96 fa ff ff
66 0f 1f 44 00 00
```

## CPU 1

```
<schedule>:
55
48 89 e5
→ 0f 1f 2e 00 00 错误指令
5d
48 8b 04 25 80 d0 15
81
48 8b 00
e9 96 fa ff ff
66 0f 1f 44 00 00
```

0f 1f 44 00 00 nop

e8 37 2e 00 00 callq ffffffff810f7430 <ftrace\_caller>

0f 1f 2e 00???

0f 1f 2e 00 00

```
0f 1f 44 00 00 nop
e8 37 2e 00 00 callq ffffffff810f7430 <ftrace_caller>
```

# 0f 1f 2e 00???

- **BOOM!**
- **CRASH!**
- **GENERAL PROTECTION FAULT!**
- **REBOOT!** 重启

# How to go from this

<schedule>:

55	push	%rbp
48 89 e5	mov	%rsp,%rbp
0f 1f 44 00 00	nop	
5d	pop	%rbp
48 8b 04 25 80 d0 15	mov	0xffffffff8115d080,%rax
81		
48 8b 00	mov	(%rax),%rax
e9 96 fa ff ff	jmpq	ffffffff810f40a0 <__schedule>
66 0f 1f 44 00 00	nopw	0x0(%rax,%rax,1)

0f 1f 44 00 00 nop

e8 37 2e 00 00 callq ffffffff810f7430 <ftrace\_caller>

# to this?

```
<schedule>:
  55                push    %rbp
  48 89 e5          mov     %rsp,%rbp
  e8 37 2e 00 00    callq   ffffffff810f7430 <ftrace_caller>
  5d                pop     %rbp
  48 8b 04 25 80 d0 15 mov     0xffffffff8115d080,%rax
  81
  48 8b 00          mov     (%rax),%rax
  e9 96 fa ff ff    jmpq    ffffffff810f40a0 <__schedule>
  66 0f 1f 44 00 00 nopw    0x0(%rax,%rax,1)
```

0f 1f 44 00 00 nop

e8 37 2e 00 00 callq ffffffff810f7430 <ftrace\_caller>



# Breakpoints

使用断点解决这个问题

2021年5月12日 rtoax



# Breakpoints

<schedule>:

55	push	%rbp
48 89 e5	mov	%rsp,%rbp
0f 1f 44 00 00	nop	
5d	pop	%rbp
48 8b 04 25 80 d0 15	mov	0xffffffff8115d080,%rax
81		
48 8b 00	mov	(%rax),%rax
e9 96 fa ff ff	jmpq	ffffffff810f40a0 <__schedule>
66 0f 1f 44 00 00	nopw	0x0(%rax,%rax,1)

0f 1f 44 00 00 nop

cc 1f 44 00 00 <bp>nop

cc 37 2e 00 00 <bp>callq ffffffff810f7430 <ftrace\_caller>

e8 37 2e 00 00 callq ffffffff810f7430 <ftrace\_caller>

# Breakpoints

<schedule>:

55		push	%rbp
48 89 e5		mov	%rsp,%rbp
cc 1f 44 00 00	<bp>	nop	
5d		pop	%rbp
48 8b 04 25 80 d0 15		mov	0xffffffff8115d080,%rax
81			
48 8b 00		mov	(%rax),%rax
e9 96 fa ff ff		jmpq	ffffffff810f40a0 <__schedule>
66 0f 1f 44 00 00		nopw	0x0(%rax,%rax,1)

0f 1f 44 00 00 nop

cc 1f 44 00 00 <bp>nop

cc 37 2e 00 00 <bp>callq ffffffff810f7430 <ftrace\_caller>

e8 37 2e 00 00 callq ffffffff810f7430 <ftrace\_caller>

# Breakpoints

<schedule>:

```
55          push    %rbp
48 89 e5     mov     %rsp,%rbp
cc 37 2e 00 00 <bp>callq  ffffffff810f7430 <ftrace_caller>
5d          pop     %rbp
48 8b 04 25 80 d0 15 mov     0xffffffff8115d080,%rax
81
48 8b 00     mov     (%rax),%rax
e9 96 fa ff ff jmpq    ffffffff810f40a0 <__schedule>
66 0f 1f 44 00 00 nopw    0x0(%rax,%rax,1)
```

0f 1f 44 00 00 nop

cc 1f 44 00 00 <bp>nop

cc 37 2e 00 00 <bp>callq ffffffff810f7430 <ftrace\_caller>

e8 37 2e 00 00 callq ffffffff810f7430 <ftrace\_caller>

# Breakpoints

<schedule>:

55		push	%rbp
48 89 e5		mov	%rsp,%rbp
e8 37 2e 00 00		callq	ffffffff810f7430 <ftrace_caller>
5d		pop	%rbp
48 8b 04 25 80 d0 15		mov	0xffffffff8115d080,%rax
81			
48 8b 00		mov	(%rax),%rax
e9 96 fa ff ff		jmpq	ffffffff810f40a0 <__schedule>
66 0f 1f 44 00 00		nopw	0x0(%rax,%rax,1)

0f 1f 44 00 00 nop

cc 1f 44 00 00 <bp>nop

cc 37 2e 00 00 <bp>callq fffffffff810f7430 <ftrace\_caller>

e8 37 2e 00 00 callq fffffffff810f7430 <ftrace\_caller>

# Registering with Ftrace

- Call to `register_ftrace_function()`
- Requires an `ftrace_ops` descriptor
- Static `ftrace_ops`
  - function and function\_graph
  - function probes (schedule:traceoff)
  - stack tracer
  - latency tracers
- Dynamic `ftrace_ops`
  - perf
  - kprobes

例:

```
//kernel/kprobes.c
enable_kprobe
arm_kprobe
    arm_kprobe_ftrace
    __arm_kprobe_ftrace
    register_ftrace_function
```



# ftrace\_ops

```
struct ftrace_ops {  
    ftrace_func_t          func;  
    struct ftrace_ops      *next;  
    unsigned long          flags;  
    int __percpu            *disabled;  
    void                   *private;  
#ifdef CONFIG_DYNAMIC_FTRACE  
    struct ftrace_hash      *notrace_hash;  
    struct ftrace_hash      *filter_hash;  
    struct mutex            regex_lock;  
#endif  
};
```

# ftrace\_ops.flags

- **ENABLED** 使能
  - set by ftrace, when ops is recording
- **DYNAMIC**
  - set by ftrace when ops is dynamically allocated
- **CONTROL**
  - set by perf

```
enum {  
    FTRACE_OPS_FL_ENABLED           = BIT(0),  
    FTRACE_OPS_FL_DYNAMIC           = BIT(1),  
    FTRACE_OPS_FL_SAVE_REGS         = BIT(2),  
    FTRACE_OPS_FL_SAVE_REGS_IF_SUPPORTED = BIT(3),  
    FTRACE_OPS_FL_RECURSION_SAFE    = BIT(4),  
    FTRACE_OPS_FL_STUB               = BIT(5),  
    FTRACE_OPS_FL_INITIALIZED        = BIT(6),  
    FTRACE_OPS_FL_DELETED            = BIT(7),  
    FTRACE_OPS_FL_ADDING             = BIT(8),  
    FTRACE_OPS_FL_REMOVING           = BIT(9),  
    FTRACE_OPS_FL_MODIFYING          = BIT(10),  
    FTRACE_OPS_FL_ALLOC_TRAMP        = BIT(11),  
    FTRACE_OPS_FL_IPMODIFY           = BIT(12),  
    FTRACE_OPS_FL_PID                = BIT(13),  
    FTRACE_OPS_FL_RCU                = BIT(14),  
    FTRACE_OPS_FL_TRACE_ARRAY        = BIT(15),  
    FTRACE_OPS_FL_PERMANENT          = BIT(16),  
    FTRACE_OPS_FL_DIRECT             = BIT(17),  
};
```

# ftrace\_ops.flags

- **SAVE\_REGS**
  - set by caller, to record regs
  - fails if saving regs is not supported
- **SAVE\_REGS\_IF\_SUPPORTED**
  - set by caller, save regs if supported
  - doesn't fail register if not supported
- **RECURSION\_SAFE** 递归
  - If ftrace\_ops.func handles recursion
  - Otherwise, ftrace will handle it

# ftrace\_ops.flags

- STUB 存根
  - used by ftrace for stub functions
- INITIALIZED
  - used by ftrace when ftrace\_ops is first used
- DELETED
  - ftrace\_ops has been deleted
  - used by ftrace buffer instances

# ftrace\_ops hashes

- regex\_lock
  - used to protect the hashes
- notrace\_hash
  - what functions not to trace 不tracing的函数
  - empty means OK to trace all
- filter\_hash
  - what functions to trace
  - empty means to trace all
- Functions in notrace\_hash will not be traced even if they exist in filter\_hash

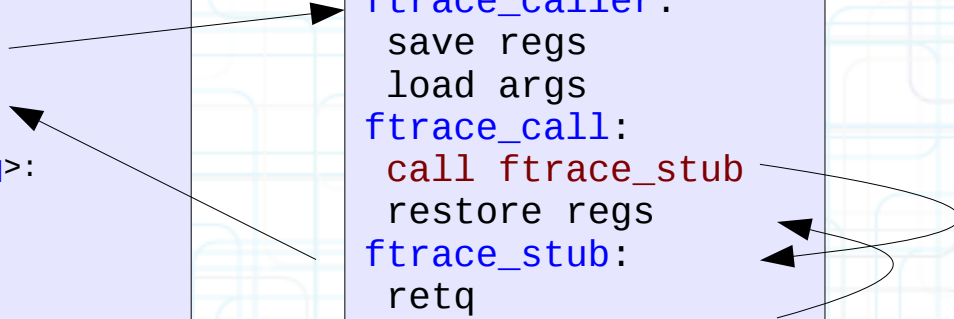


蹦床

# ftrace\_caller trampoline

```
<schedule>:  
push    %rbp  
mov     %rsp,%rbp  
call    ftrace_caller  
pop     %rbp  
[...]  
<preempt_schedule_irq>:  
push    %rbp  
mov     %rsp,%rbp  
push    %rbx  
nop  
pop     %rbp  
[...]  
<_cond_resched>:  
push    %rbp  
mov     %rsp,%rbp  
push    %rbx  
nop  
pop     %rbp  
[...]  
<yield>:  
push    %rbp  
mov     %rsp,%rbp  
push    %rbx  
call    ftrace_regs_caller  
pop     %rbp  
[...]
```

```
ftrace_caller:  
save regs  
load args  
ftrace_call:  
call    ftrace_stub  
restore regs  
ftrace_stub:  
retq
```





# ftrace\_caller trampoline

```
<schedule>:  
push    %rbp  
mov     %rsp,%rbp  
call ftrace_caller  
pop     %rbp  
[...]  
<preempt_schedule_irq>:  
push    %rbp  
mov     %rsp,%rbp  
push    %rbx  
nop  
pop     %rbp  
[...]  
<_cond_resched>:  
push    %rbp  
mov     %rsp,%rbp  
push    %rbx  
nop  
pop     %rbp  
[...]  
<yield>:  
push    %rbp  
mov     %rsp,%rbp  
push    %rbx  
call ftrace_regs_caller  
pop     %rbp  
[...]
```

```
ftrace_caller:  
save regs  
load args  
ftrace_call:  
call func_trace  
restore regs  
ftrace_stub:  
retq
```

```
void func_trace()  
{  
    /* trace */  
}
```

蹦床

# ftrace\_caller trampoline

```
<schedule>:  
push    %rbp  
mov     %rsp,%rbp  
call ftrace_caller  
pop     %rbp  
[...]  
<preempt_schedule_irq>:  
push    %rbp  
mov     %rsp,%rbp  
push    %rbx  
nop  
pop     %rbp  
[...]  
<_cond_resched>:  
push    %rbp  
mov     %rsp,%rbp  
push    %rbx  
nop  
pop     %rbp  
[...]  
<yield>:  
push    %rbp  
mov     %rsp,%rbp  
push    %rbx  
call ftrace_regs_caller  
pop     %rbp  
[...]
```

```
ftrace_caller:  
save regs  
load args  
ftrace_call:  
call func_trace  
restore regs  
ftrace_stub:  
retq
```

```
void func_trace()  
{  
    /* trace */  
}
```

ftrace\_ops.func

多个回调

# Multiple callbacks?

- Direct call works fine
- Multiple calls requires a list operation
- All functions being traced will call the list function

多个回调函数

# Multiple callbacks?

```
<schedule>:
push    %rbp
mov     %rsp,%rbp
call ftrace_caller
pop     %rbp
[...]
<preempt_schedule_irq>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
nop
pop     %rbp
[...]
<_cond_resched>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
nop
pop     %rbp
[...]
<yield>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
call ftrace_regs_caller
pop     %rbp
[...]
```

ftrace\_caller:

```
save regs
load args
ftrace_call:
call list_func
restore regs
ftrace_stub:
retq
```

void list\_func()

```
{
    /* iterate */ 迭代遍历
}
```

void func1\_trace()

```
{
    /* trace */
}
```

void func2\_trace()

```
{
    /* trace */
}
```

# Example

`function tracer->all functions`

`perf->schedule`

- Run function tracer on all functions
- Run perf on just the scheduler

function tracer->all functions  
perf-> schedule

# Example

Want to trace  
do\_IRQ()?

list\_func()

Nope

perf

Yep!

function\_tracer



function tracer->all functions  
perf-> schedule

# Example

Want to trace  
`raw_spin_lock()`?

`list_func()`

Nope

perf

Yep!

function\_tracer

function tracer->all functions  
perf-> schedule

# Example

Want to trace  
rcu\_read\_lock()?

list\_func()

Nope

perf

Yep!

function\_tracer

function tracer->all functions  
perf-> schedule

# Example

Want to trace  
rcu\_read\_unlock()?

list\_func()

Nope

perf

Yep!

function\_tracer

function tracer->all functions

perf-> schedule

# Example

Want to trace  
schedule()?

list\_func()

Yep!

perf

Yep!

function\_tracer

function tracer->all functions  
perf-> schedule

# Multiple callbacks?

调度函数

跳板

```
<schedule>:  
push    %rbp  
mov     %rsp,%rbp  
call    ftrace_caller  
pop     %rbp  
[...]  
<preempt_schedule_irq>:  
push    %rbp  
mov     %rsp,%rbp  
push    %rbx  
call    ftrace_caller  
pop     %rbp  
[...]  
<_cond_resched>:  
push    %rbp  
mov     %rsp,%rbp  
push    %rbx  
call    ftrace_caller  
pop     %rbp  
[...]  
<yield>:  
push    %rbp  
mov     %rsp,%rbp  
push    %rbx  
call    ftrace_caller  
pop     %rbp  
[...]
```

```
ftrace_caller:  
save regs  
load args  
ftrace_call:  
call    list_func  
restore regs  
ftrace_stub:  
retq
```

多回调 (禁止抢占)

```
void list_func()  
{  
    /* iterate and  
    check hash of ops */  
}
```

函数trace

```
void function_trace()  
{  
    /* function  
    tracing */  
}
```

perf trace

```
void perf_func()  
{  
    /* perf  
    profiling */  
}
```

# Multiple callbacks?

```
<schedule>:  
push    %rbp  
mov     %rsp,%rbp  
call    ftrace_caller  
pop     %rbp  
[...]  
<preempt_schedule_irq>:  
push    %rbp  
mov     %rsp,%rbp  
push    %rbx  
call    ftrace_caller  
pop     %rbp  
[...]  
<_cond_resched>:  
push    %rbp  
mov     %rsp,%rbp  
push    %rbx  
call    ftrace_caller  
pop     %rbp  
[...]  
<yield>:  
push    %rbp  
mov     %rsp,%rbp  
push    %rbx  
call    ftrace_caller  
pop     %rbp  
[...]
```

```
ftrace_caller:  
save regs  
load args  
ftrace_call:  
call    list_func  
restore regs  
ftrace_stub:  
retq
```

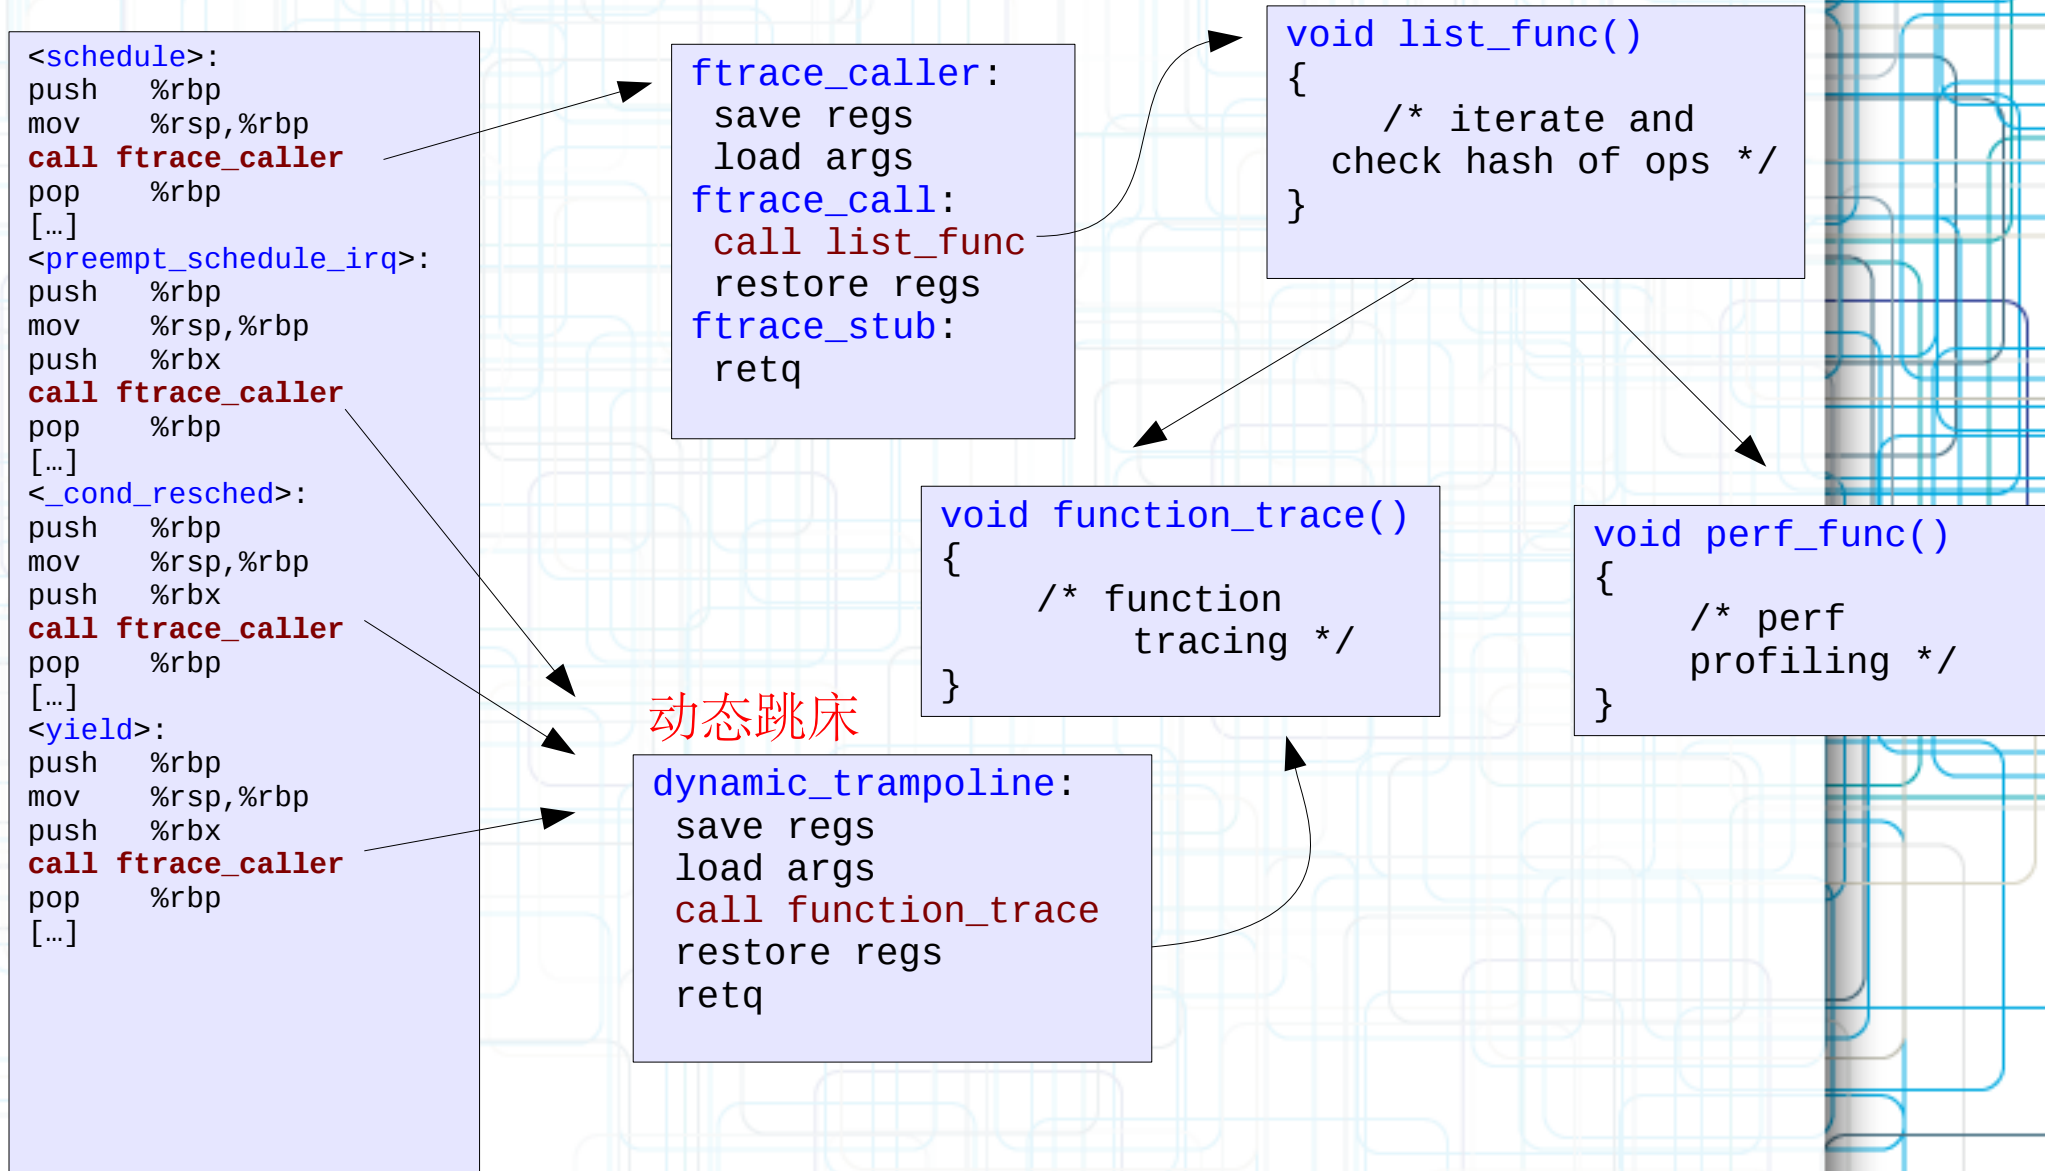
```
void list_func()  
{  
    /* iterate and  
    check hash of ops */  
}
```

```
void function_trace()  
{  
    /* function  
    tracing */  
}
```

```
void perf_func()  
{  
    /* perf  
    profiling */  
}
```

动态跳床

```
dynamic_trampoline:  
save regs  
load args  
call    function_trace  
restore regs  
retq
```





# Problems with Dynamic Trampolines

- When can you free them?
- When are they not in use?
- Would RCU help?

动态跳床的问题

1. 什么时候free?
2. 什么时候不被使用?
3. RCU有帮助吗?

# Dynamic trampolines


Task



```
dynamic_trampoline:  
    save regs  
    load args  
    call function_trace  
    restore regs  
    retq
```

# Dynamic trampolines

Task



Preempted!  
被抢占

```
dynamic_trampoline:  
    save regs  
    load args  
    call function_trace  
    restore regs  
    retq
```

# What about dynamic ftrace\_ops

动态和静态 ftrace\_ops 是有区别的

- Remember, there's a difference between dynamic and static ftrace\_ops

**ftrace** 检测 ftrace\_ops 动态与否

- Ftrace detects ftrace\_ops that are dynamic

总是使用 list function

- Always uses the list function
  - it disables preemption 禁止抢占
  - and is static **list** function 是静态的

# Dynamic ftrace\_ops

```
<schedule>:
push    %rbp
mov     %rsp,%rbp
call ftrace_caller
pop     %rbp
[...]
<preempt_schedule_irq>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
nop
pop     %rbp
[...]
<_cond_resched>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
nop
pop     %rbp
[...]
<yield>:
push    %rbp
mov     %rsp,%rbp
push    %rbx
call ftrace_regs_caller
pop     %rbp
[...]
```

```
ftrace_caller:
save regs
load args
ftrace_call:
call list_func
restore regs
ftrace_stub:
retq
```

静态的，禁止抢占的

```
void list_func()
{
    preempt_disable_notraced();
    /* iterate */
    preempt_enable_notraced();
}
```

```
void dynamic_ops_func()
{
    /* trace */
}
```

```
#define preempt_disable_notraced() \
do { \
    __preempt_count_inc(); \
    barrier(); \
} while (0)
```

# Knowing when to free

- If there was a way to know no more tasks were on the trampoline or function
- There will be a way in coming 3.18

当有一种方法可以知道没有tasks在蹦床或者函数时；



# call\_rcu\_tasks()

- Call a function after all tasks
  - have voluntarily scheduled 自愿调度
  - in userspace 用户空间
  - are idle **idle**任务

动态蹦床

# Dynamic trampolines

```
<schedule>:  
push    %rbp  
mov     %rsp,%rbp  
call dynamic_tramp  
pop     %rbp  
[...]
```

Task

```
dynamic_trampoline:  
save regs  
load args  
call function_trace  
restore regs  
retq
```

# Dynamic trampolines

```
<schedule>:  
push    %rbp  
mov     %rsp,%rbp  
call dynamic_tramp  
pop     %rbp  
[...]
```

tramp: 流浪汉

```
dynamic_trampoline:  
save regs  
load args  
call function_trace  
restore regs  
retq
```

Task

Preempted!  
抢占

# Dynamic trampolines

```
<schedule>:  
push    %rbp  
mov     %rsp,%rbp  
call ftrace_stub stub: 存根  
pop     %rbp  
[...]
```

Task

Preempted!

```
dynamic_trampoline:  
save regs  
load args  
call function_trace  
restore regs  
retq
```

# Dynamic trampolines

```
<schedule>:  
push    %rbp  
mov     %rsp,%rbp  
call ftrace_stub  
pop     %rbp  
[...]
```

Task

Preempted!

```
dynamic_trampoline:  
save regs  
load args  
call function_trace  
restore regs  
retq
```

**call\_rcu\_task()**

# Dynamic trampolines

```
<schedule>:  
push    %rbp  
mov     %rsp,%rbp  
call ftrace_stub  
pop     %rbp  
[...]
```

```
dynamic_trampoline:  
save regs  
load args  
call function_trace  
restore regs  
retq
```

Task

- Voluntary schedule
- In idle
- In userspace

**call\_rcu\_task()**



# Dynamic trampolines

```
<schedule>:  
push    %rbp  
mov     %rsp,%rbp  
call ftrace_stub  
pop     %rbp  
[...]
```

Task

- Voluntary schedule
- In idle
- In userspace

```
dynamic_trampoline:  
save regs  
load args  
call function_trace  
restore regs  
retq
```

# fentry

**mcount**不能记录参数

- mcount can't record parameters
- New feature of gcc
  - starting with gcc 4.6.0
    - Added by Andi Kleen
  - for x86\_64 only (for now)
- gcc -pg -mfentry **-mfentry**

<schedule>:

55	push	%rbp
48 89 e5	mov	%rsp,%rbp
e8 37 2e 00 00	callq	ffffffff810f7430 <mcount>
5d	pop	%rbp
48 8b 04 25 80 d0 15	mov	0xffffffff8115d080,%rax
81		
48 8b 00	mov	(%rax),%rax
e9 96 fa ff ff	jmpq	ffffffff810f40a0 <__schedule>
66 0f 1f 44 00 00	nopw	0x0(%rax,%rax,1)

<posix\_cpu\_timer\_set>:

55	push	%rbp
48 89 e5	mov	%rsp,%rbp
41 57	push	%r15
41 56	push	%r14
41 55	push	%r13
41 54	push	%r12
53	push	%rbx
48 83 ec 30	sub	\$0x30,%rsp
e8 1a 81 0b 00	callq	ffffffff810f7430 <mcount>
48 8b 47 70	mov	0x70(%rdi),%rax
49 89 ff	mov	%rdi,%r15

<posix\_cpu\_timer\_set>:

e8 eb 3c 0b 00	callq	ffffffff810f0af0 <__fentry__>
41 57	push	%r15
41 56	push	%r14
49 89 ff	mov	%rdi,%r15
41 55	push	%r13
41 54	push	%r12
49 89 d5	mov	%rdx,%r13
55	push	%rbp
53	push	%rbx

# fentry

```
<posix_cpu_timer_set>:
```

```
call ftrace_caller
```

```
push    %r15
```

```
push    %r14
```

```
mov     %rdi,%r15
```

```
push    %r13
```

```
push    %r12
```

```
ftrace_caller:
```

```
save regs
```

```
load args
```

```
ftrace_call:
```

```
call func_trace
```

```
restore regs
```

```
ftrace_stub:
```

```
retq
```

```
void func_trace()
```

```
{
```

```
    /* trace */
```

```
}
```

# Live Kernel Patching!

<posix\_cpu\_timer\_set>:

**call ftrace\_caller**

push %r15  
push %r14  
mov %rdi,%r15  
push %r13  
push %r12

ftrace\_caller:

save regs

load args

ftrace\_call:

**call func\_trace**

restore regs

ftrace\_stub:

retq

void lkp()

{

/\* change  
return reg \*/

}

<posix\_cpu\_timer\_set>:

**nop**

push %r15  
push %r14  
mov %rdi,%r15  
push %r13  
push %r12

# Another Solution

vmlinux:

```
<schedule>:
nop
push    %rbp
mov     %rsp,%rbp
pop     %rbp
[...]
<preempt_schedule_irq>:
nop
push    %rbp
mov     %rsp,%rbp
push    %rbx
pop     %rbp
[...]
<_cond_resched>:
nop
push    %rbp
mov     %rsp,%rbp
push    %rbx
pop     %rbp
[...]
<yield>:
nop
push    %rbp
mov     %rsp,%rbp
push    %rbx
pop     %rbp
[...]
```

<ftrace\_pages>

```
ip      = 0xffffffff81087ad4
flags   = 0
ip      = 0xffffffff81087b14
flags   = 0
ip      = 0xffffffff81087bd5
flags   = 0
ip      = 0xffffffff81087c41
flags   = 0
ip      = 0xffffffff810a7aa0
flags   = 0
ip      = 0xffffffff810a7bd4
flags   = 0
ip      = 0xffffffff810a7d34
flags   = 0
ip      = 0xffffffff810a7d7d
flags   = 0
ip      = 0xffffffff810f45f4
flags   = 0
ip      = 0xffffffff810f4635
flags   = 0
ip      = 0xffffffff810f4684
flags   = 0
ip      = 0xffffffff810f4734
flags   = 0
[...]
```



# Another Solution

vmlinux:

```
<schedule>:
nop
push    %rbp
mov     %rsp,%rbp
pop     %rbp
[...]
<preempt_schedule_irq>:
nop
push    %rbp
mov     %rsp,%rbp
push    %rbx
pop     %rbp
[...]
<_cond_resched>:
nop
push    %rbp
mov     %rsp,%rbp
push    %rbx
pop     %rbp
[...]
<yield>:
nop
push    %rbp
mov     %rsp,%rbp
push    %rbx
pop     %rbp
[...]
```

```
<preempt_schedule_irq>:
nop
push    %r15
push    %r14
mov     %rdi,%r15
push    %r13
push    %r12
```

<ftrace\_pages>

```
ip      = 0xffffffff81087ad4
flags   = 0
ip      = 0xffffffff81087b14
flags   = 0
ip      = 0xffffffff81087bd5
flags   = 0
ip      = 0xffffffff81087c41
flags   = 0
ip      = 0xffffffff810a7aa0
flags   = 0
ip      = 0xffffffff810a7bd4
flags   = 0
ip      = 0xffffffff810a7d34
flags   = 0
ip      = 0xffffffff810a7d7d
flags   = 0
ip      = 0xffffffff810f45f4
flags   = 0
ip      = 0xffffffff810f4635
flags   = 0
ip      = 0xffffffff810f4684
flags   = 0
ip      = 0xffffffff810f4734
flags   = 0
[...]
```

# Another Solution

vmlinux:

```
<schedule>:
nop
push    %rbp
mov     %rsp,%rbp
pop     %rbp
[...]
<preempt_schedule_irq>:
nop
push    %rbp
mov     %rsp,%rbp
push    %rbx
pop     %rbp
[...]
<_cond_resched>:
nop
push    %rbp
mov     %rsp,%rbp
push    %rbx
pop     %rbp
[...]
<yield>:
nop
push    %rbp
mov     %rsp,%rbp
push    %rbx
pop     %rbp
[...]
```

```
<preempt_schedule_irq>:
nop
push    %r15
push    %r14
mov     %rdi,%r15
push    %r13
push    %r12
```

<ftrace\_pages>

```
ip      = 0xffffffff81087ad4
flags   = 0
ip      = 0xffffffff81087b14
flags   = 0
ip      = 0xffffffff81087bd5
flags   = 0
ip      = 0xfffffffffa0014466
flags   = 0
ip      = 0xffffffff810a7aa0
flags   = 0
ip      = 0xffffffff810a7bd4
flags   = 0
ip      = 0xffffffff810a7d34
flags   = 0
ip      = 0xffffffff810a7d7d
flags   = 0
ip      = 0xffffffff810f45f4
flags   = 0
ip      = 0xffffffff810f4635
flags   = 0
ip      = 0xffffffff810f4684
flags   = 0
ip      = 0xffffffff810f4734
flags   = 0
[...]
```

# Another Solution

vmlinux:

```
<schedule>:
nop
push    %rbp
mov     %rsp,%rbp
pop     %rbp
[...]
<preempt_schedule_irq>:
jmp preempt_sched2
push    %rbp
mov     %rsp,%rbp
push    %rbx
pop     %rbp
[...]
<_cond_resched>:
nop
push    %rbp
mov     %rsp,%rbp
push    %rbx
pop     %rbp
[...]
<yield>:
nop
push    %rbp
mov     %rsp,%rbp
push    %rbx
pop     %rbp
[...]
```

```
<preempt_schedule_irq>:
nop
push    %r15
push    %r14
mov     %rdi,%r15
push    %r13
push    %r12
```

<ftrace\_pages>

```
ip      = 0xffffffff81087ad4
flags   = 0
ip      = 0xffffffff81087b14
flags   = 0
ip      = 0xffffffff81087bd5
flags   = 0
ip      = 0xfffffffffa0014466
flags   = 0
ip      = 0xffffffff810a7aa0
flags   = 0
ip      = 0xffffffff810a7bd4
flags   = 0
ip      = 0xffffffff810a7d34
flags   = 0
ip      = 0xffffffff810a7d7d
flags   = 0
ip      = 0xffffffff810f45f4
flags   = 0
ip      = 0xffffffff810f4635
flags   = 0
ip      = 0xffffffff810f4684
flags   = 0
ip      = 0xffffffff810f4734
flags   = 0
[...]
```

# Instead of this

```
<posix_cpu_timer_set>:
```

```
call ftrace_caller
```

```
push    %r15  
push    %r14  
mov     %rdi,%r15  
push    %r13  
push    %r12
```

```
ftrace_caller:
```

```
save regs
```

```
load args
```

```
ftrace_call:
```

```
call func_trace
```

```
restore regs
```

```
ftrace_stub:
```

```
retq
```

```
void lkp()
```

```
{
```

```
    /* change  
    return reg */
```

```
}
```

```
<posix_cpu_timer_set>:
```

```
nop
```

```
push    %r15  
push    %r14  
mov     %rdi,%r15  
push    %r13  
push    %r12
```

# Have this!

```
<posix_cpu_timer_set>:  
jmp posix_cpu2  
push    %r15  
push    %r14  
mov     %rdi,%r15  
push    %r13  
push    %r12
```

```
<posix_cpu_timer_set>:  
call ftrace_caller  
push    %r15  
push    %r14  
mov     %rdi,%r15  
push    %r13  
push    %r12
```

```
ftrace_caller:  
    save regs  
    load args  
ftrace_call:  
    call func_trace  
    restore regs  
ftrace_stub:  
    retq
```

# Questions?

Yeah right!  
Like we have time